

We plan to test whether changes in the University of Michigan Survey of Consumers' sentiment measures anticipate near-term macro-outcomes. The topic could be: **Do monthly shifts in Consumer Sentiment and its components (Current Conditions, Expectations) predict subsequent changes in unemployment and retail sales 1–6 months ahead, and are the relationships different across recessions (e.g., 2008 vs. 2020)?** We plan to obtain the sentiment series from the Survey of Consumers website (scraping the historical index table with rvest) and, for robustness and easy access, also pull equivalent monthly series along with outcome series such as the unemployment rate (BLS) and retail sales (Census)—all programmatically acquired from the web. The unit of analysis will be a month, with approximately 550–600 rows (late 1970s to present) and about 5 to 10 variables (sentiment levels and changes, unemployment, retail sales, recession flags, etc.; numeric and ordinal data types). We plan to use dplyr for data management, produce time-series line charts (with recession shading), lagged scatterplots of sentiment vs. future outcomes, and rolling correlations, etc. I'll also fit simple lagged OLS regressions and cross-correlation summaries to quantify lead–lag structure. The deliverables are expected to include clear visualizations that answer the question and a brief discussion of limitations.